



# Linux Containers

A software architecture lecture

Josef Karasek

Quality Engineer, Red Hat Middleware

2017-04-03

# Agenda

- Linux containers
  - Motivation
  - Containers as a packaging mechanism
  - Containers as process isolation

# When does SW product provide value?

- SW development is hard and long process
- Production SW needs to be maintained and supported


Development & Testing




# Systems have become more complex

- More subsystems involved
- More infrastructure to manage
- Development cycle getting shorter - change is very common

# Topologies used today




**Impact = N**




**Impact = 1**

450 microservices




500+ microservices



NETFLIX

500+ microservices



Source:

Netflix: <http://www.slideshare.net/BruceWong3/the-case-for-chaos>

Twitter: <https://twitter.com/adrianco/status/441883572618948608>

Hail-o: <https://sudo.hailoapp.com/services/2015/03/09/journey-into-a-microservice-world-part-3/>

# Immutable architecture

Evolving SW architecture at scale

- Systems parts are not modified but replaced by new ones
- Immutable Infrastructure brings stability & consistency to the environment (update success / fail?)
- Easy to move forward, but also backward
- Involves also process/environment change

# Trash Your Servers and Burn Your Code


-- Chad Fowler





# Linux containers

# Containers as a packaging mechanism

- Code and its runtime dependencies bundled together
- Format of container is well understood
- Commonly a .tar archive of
  - File system
  - Static binary



# Containers as a packaging mechanism II



# What's inside a container

Inside / Outside

## Code

Compiled binary  
Shared libraries  
Configuration scripts  
JRE, Python...

## Configuration

Injected at runtime - easy  
customization  
Files, environment  
variables...

## Data

Persisted outside  
Containers can be  
restarted with no data  
loss  
Data has its own lifecycle,  
independent from code

# Example container: MySQL

**Code**

`mysqld`

**Configuration**

`/etc/my.cnf`

**Data**

`/var/lib/mysql`

Runtime dependencies


```
$ ldd /usr/libexec/mysqld
linux-vdso.so.1
libsystemd.so.0 => /lib64/libsystemd.so.0
libpthread.so.0 => /lib64/libpthread.so.0
...
```

# Containers as process isolation

A container is just a fancy process

Namespaces - restrict what resources the container can use


CGroups - define how much of a resource the container can use



# Namespaces

Provide containers with their own view of the underlying Linux system

- pid
- mnt
- net
- ipc (inter-process comm)
- uts (hostname, domainname)
- user



# Control Groups

## Resource control and accounting

- Cpu share
- Cpuset
- Memory allocation
  - Soft vs. hard limits
- I/O
- Devices cgroup
- Freezer group
- Accounting (memory page  $\sim$  4kB)

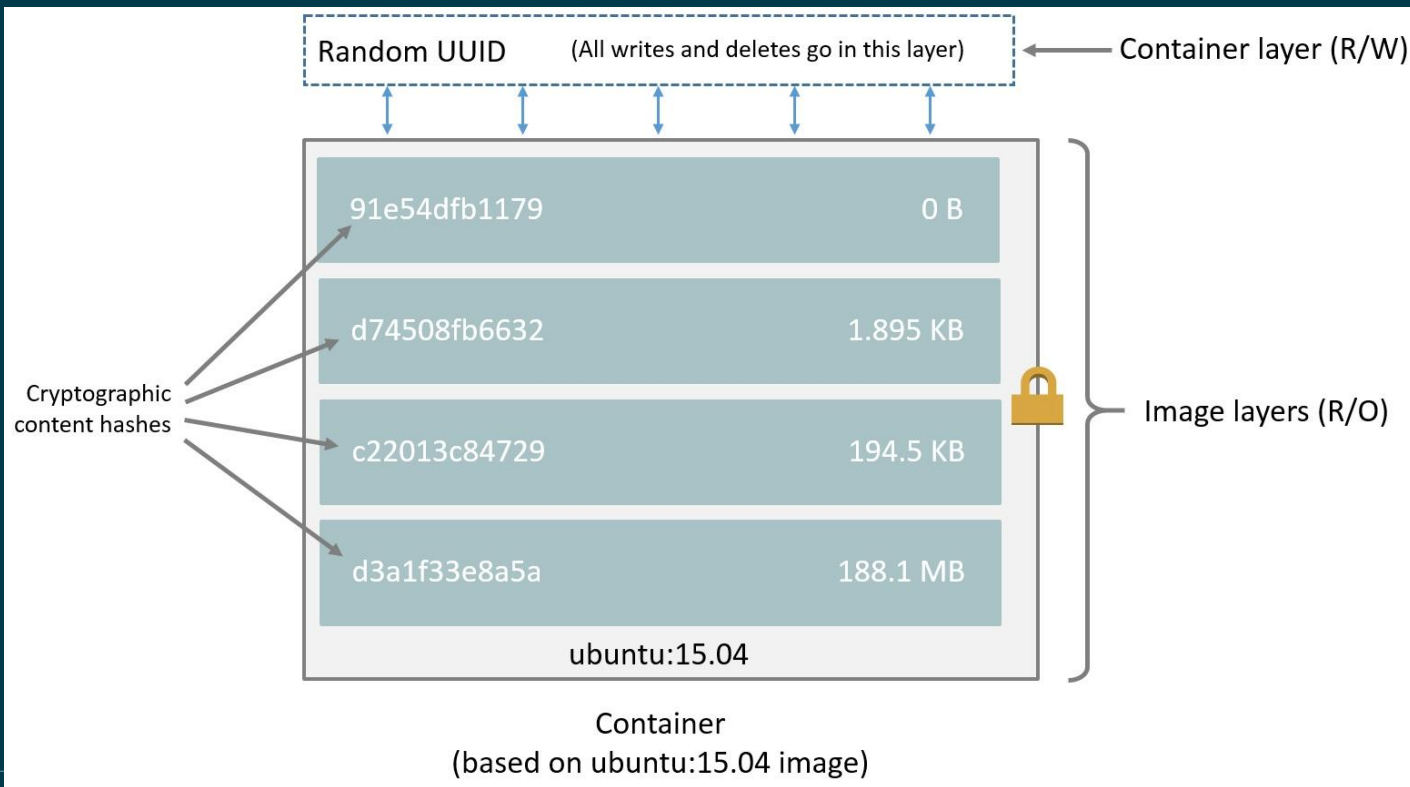


# Union filesystem & Copy on Write

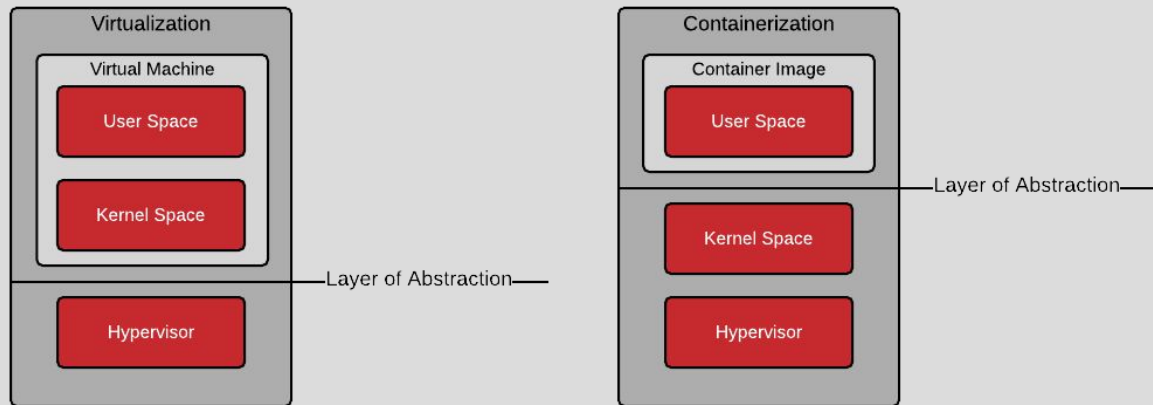
Killer feature for usable container

- Union filesystem
  - Image is a set of layers
  - Reusing/combining layers is efficient
- Copy on write (CoW)
  - Container is an image and a thin writable layer (container layer)
  - Fast spawning/deleting container
  - Deleting container = deleting only a layer

# Union filesystem & Copy on Write II



# Containers are different from virtualization



# Docker container image

Command:

```
docker pull registry.access.redhat.com/rhel7/rhel:latest
```

Decomposition:

access.registry.redhat.com

/

rhel7

/

rhel

:

latest

Generalization:

Registry Server

/

namespace

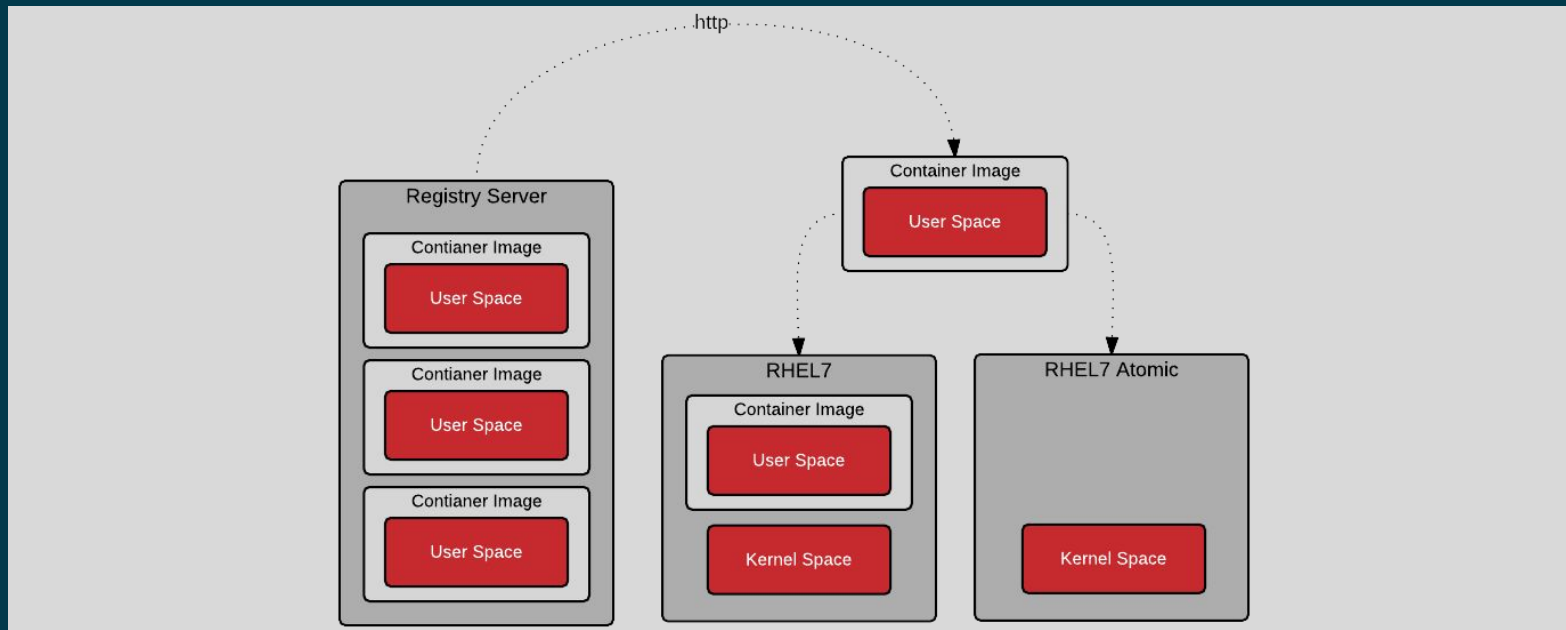
/

repo

:

tag

# Registry infrastructure



# Demo time!





# THANK YOU



[plus.google.com/+RedHat](https://plus.google.com/+RedHat)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[twitter.com/RedHatNews](https://twitter.com/RedHatNews)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)